

Where Did the Agent Go Wrong? Evaluating Failure Attribution Against Counterfactual Rescue

Anonymous ACL submission

Abstract

Self-improving agent loops start by attributing each failure to a step, usually by asking an LLM to read the trace. We evaluate attribution against counterfactual rescue: for every step we measure how much one oracle correction raises the success of continued replays above uncorrected replays from that step, and score a method by this rescue gain at the step it names. On 300 natural failures from a self-improvement loop in τ^2 -bench retail and airline and in AppWorld, with registered analyses, a zero-cost rule naming the first state-changing action scored highest under the registered prompts (0.42, against 0.20–0.34), not significantly above the best LLM method; judges asked post hoc for the measured quantity came within 0.02. The rule’s lead rests on an oracle that reads the grading: it came from failures best rescued at a write, and an oracle shown only the failed episode rarely rescued writes and ranked the rule seventh of ten. Without null replays, early steps earn credit for chance recoveries and the ranking reverses. Counterfactual search did not detectably beat one judge call at seven times the cost when only search lacked the grading. In a repair experiment, neither these attributions nor the step of largest rescue gain, even with a diagnosis and correction that kept most of their gain on fresh replays, brought a detectable gain on held-out tasks over the loop’s own analysis; the experiment cannot exclude gains the size of one accepted round.

1 Introduction

Agents that improve themselves from their own failures share a first step. An analyst model reads a failed episode, decides where it went wrong, and hands that diagnosis to a proposer that edits the agent’s prompts, tools or code (Shinn et al., 2023; Zhao et al., 2024; Zhang et al., 2026b; Chen et al., 2026; Xia et al., 2026). If the diagnosis points at the wrong step, the proposer may repair the wrong thing. Failure attribution has therefore become a

task of its own, with benchmarks of human-labelled multi-agent logs (Zhang et al., 2025; Liu et al., 2026), taxonomies of failure modes (Cemri et al., 2025, 2026; Zhu et al., 2025), and methods that replay a failure with a step changed (Zhang et al., 2026a; Ma et al., 2025; Lin et al., 2026; Bonagiri et al., 2026; Shah, 2026).

Three questions are rarely asked together: is an attribution correct in the sense a repair loop needs, is it worth its cost, and does a more accurate attribution make the loop’s repairs better? A human label says which step an annotator blames; a repair loop needs a step whose correction would have made the episode succeed. We evaluate replay-based methods and LLM judges against a reference defined by rescue, report cost next to the score, and hand each attribution to a working loop to measure the repairs it produces.

Defining that reference was the hard part: the *earliest* step whose single correction beats uncorrected replays was unstable, because a failure usually has several rescuing steps. We therefore measure a *rescue profile*: for every step, the mean gain in success over null replays across four independent corrections. A method is scored by the rescue gain at the step it names, and a failure is *rescuable* when some step reaches a gain of 0.5. The null replays, which continue the unchanged episode from the same step, matter most: without them an LLM judge would rank first and the rule that wins with them sixth.

We apply the profile to 300 natural failures of agents built by a self-improvement loop, RRSI (Xia et al., 2026), in τ^2 -bench’s retail and airline domains (Barres et al., 2025) and in AppWorld with a code agent (Trivedi et al., 2024), and compare ten attribution methods, from LLM judges of different cost and design (Zhang et al., 2025) to counterfactual search with up to 40 replays and two structural rules. Methods, metrics and analyses were registered before any rescue profile was computed. A

second registered experiment re-runs one round of RRSI at eight of its own states in τ^2 with each attribution, or none, given to the proposer, and deploys every resulting repair on 240 held-out episodes (§6); post hoc, we repeat it with a proposer that cannot read the failing traces. Our contributions:

- **A reference by rescue profile** with null controls, and what it depends on. Post hoc, the method ranking held across oracle samples, thresholds, a second oracle, the observed action as control and a screen of the corrections, but two choices change it: without null replays it reverses, because uncorrected replays from early steps often succeed by chance, and without the grading the oracle rarely rescues writes (§5).
- **A cost-aware comparison of ten attribution methods** on 300 natural failures. Under the registered reference a free structural rule was competitive with all of them: the first write scored highest under the registered prompts, not significantly above binary search ($+0.08 [-0.02, +0.18]$), and judges asked post hoc for the quantity the metric measures came within 0.02. Its lead came from writes, which the oracle repairs with the grading. Counterfactual search did not detectably beat one judge call at seven times the cost when only search lacked the grading, and at matched information led by at most 0.05 (§5).
- **A test on repairs.** Repairs drafted from first-write or binary-search steps, from the step of largest rescue gain, or from that step’s diagnosis and correction given to a proposer that could not read the traces, showed no detectable benefit on held-out tasks over those drafted from RRSI’s own analysis, for candidates or for accepted rounds. The intervals exclude differences beyond about ± 2.5 points, not the gain of one accepted round, and no group gained on average (§6).

2 Related Work

Failure attribution. Who&When introduced automated attribution of the responsible agent and step in multi-agent logs, with all-at-once, step-by-step and binary-search judges, and found step-level accuracy far below agent-level accuracy (Zhang et al., 2025); Who&When Pro revisits these labels (Liu et al., 2026). Taxonomies organise fail-

ures into modes (Cemri et al., 2025, 2026; Zhu et al., 2025, 2026); ErrorProbe learns diagnosis from its history (Li et al., 2026). MP-Bench argues that a failure often admits several plausible attributions and finds that earlier conclusions about LLMs’ weakness at attribution owe much to benchmark design (In et al., 2026). These evaluations score agreement with human or injected labels; we score the effect of correcting the named step, and likewise find that how judges are asked changes their score.

Counterfactual attribution. AgenTracer labels failures by replaying corrected steps of injected faults and trains an attributor on them (Zhang et al., 2026a); DoVer and REFLECT intervene on a trace to validate hypotheses (Ma et al., 2025; Lin et al., 2026); CausalFlow and Causal Agent Replay attribute failures by counterfactual replay (Bonagiri et al., 2026; Shah, 2026). Our reference uses the same intervention with null replays from the same step and a full profile over steps instead of the first step that flips. The profile is a graded, sampled version of but-for causation (Halpern, 2016; Pearl, 2009).

Self-improving agents. Reflexion, ExpEL and ACE turn failures into lessons or playbooks (Shinn et al., 2023; Zhao et al., 2024; Zhang et al., 2026b); HarnessFix diagnoses and repairs harness flaws (Chen et al., 2026); RRSI digests failed traces, proposes harness edits and keeps them by re-evaluation (Xia et al., 2026). Prediction of failure does not imply that intervening helps (Vasudev et al., 2026). We evaluate the diagnostic step these loops share, on failures one of them produced.

3 A Reference: Rescue Profiles

Replays. A failed episode is a sequence of agent steps $0, \dots, n-1$. A *replay from step k* keeps steps $0, \dots, k-1$ and the environment state they produced, and lets the same agent and user simulator continue; a *corrected* replay forces a given action at step k first. The agent runs at temperature 0, but replays still vary, so each outcome is a sample.

Oracle corrections. For each step an oracle (DeepSeek-V4-Pro, thinking off) sees the domain policy, the tools, the episode up to the step and the task’s grading criteria: τ^2 ’s gold actions and required information, or AppWorld’s unit tests with their expected values. It either judges the step correct or writes the action the agent should have taken,

using only what the agent knew at that point: the criteria tell it the right outcome, not identifiers or facts the agent had not yet seen. The reference is thus reparability by an oracle that knows the intended outcome, privileged information an offline evaluation may use; it does not certify that each correction was justified by what the agent had seen.

Protocol 3. It replaced an earliest-step protocol whose independent runs agreed on 42% of pilot failures (Appendix A). At every step k we draw $K = 4$ oracle samples. Each sample that corrects the step gets 2 corrected replays, and the step gets 4 null replays. The rescue gain of step k is

$$R_k = \frac{1}{K} \sum_{i=1}^K c_{k,i} (\hat{p}_{k,i}^{\text{corr}} - \hat{p}_k^{\text{null}}), \quad (1)$$

where $c_{k,i} = 1$ if sample i gives an action the replay can apply, and $c_{k,i} = 0$ if it judges the step correct or its action cannot be applied. R_k is thus the share of samples that propose a correction times the mean gain of those corrections. A null replay lets the agent act again at step k . It reproduced the observed tool calls in 91–95% of τ^2 null replays and the observed code in 58% of AppWorld ones (post hoc), so R_k compares forcing the oracle’s action with letting the agent act again, which in τ^2 is close to keeping the observed action. A failure is rescuable if $\max_k R_k \geq 0.5$, and then its best rescue step is the step with the largest gain (earliest on ties). We also record the earliest step with $R_k \geq 0.5$. Appendix Figure 1 shows one profile.

Scoring a method. A method names one step $\hat{k}$. Its score on a failure is $R(\hat{k})$, the rescue gain at that step, and 0 for an invalid step. The primary metric is the mean of $R(\hat{k})$ over rescuable failures; it rewards naming any step whose correction helps, in proportion to how much it helps. Secondary metrics include exact agreement with the best rescue step (Appendix Table 4).

4 Setup

Failures. All failures come from RRSI (Xia et al., 2026) run as a self-improvement loop on a DeepSeek-V4-Flash agent: from the harnesses it accepted over 20 rounds, in their evaluations on training tasks and in their deployment on held-out tasks. In τ^2 -bench (Barres et al., 2025) the agent serves a simulated customer through tool calls; in AppWorld (Trivedi et al., 2024) it writes Python against the APIs of nine apps and submits with a

final call. We drew 100 failed episodes per domain with a registered rule: no harness error, at most two per task and harness, tasks taken round-robin in a seeded order. They cover 36 (retail), 24 (airline) and 47 (AppWorld) tasks and average 11.7, 10.3 and 19.6 agent steps.

Methods. Every LLM method except counterfactual search reads the failure as RRSI’s analyst reads it: the task (in τ^2 with the simulated customer’s hidden instructions), the episode with numbered steps, and the grading report the agent never sees, which lists τ^2 ’s gold actions and missing information or AppWorld’s failed unit tests with their assertions. The judges therefore have the information the oracle has. Each method names one step.

- **Rules:** the *first write*, the first step that calls a state-changing tool or API (else the last step), and the *last step*.
- **All-at-once** judge (Zhang et al., 2025): one call over the whole episode, with DeepSeek-V4-Flash, V4-Pro, and V4-Pro with thinking.
- **Step-by-step** judge: one call per step, in order, until a step is judged decisive.
- **Binary search** judge: halves the episode $\lceil \log_2 n \rceil$ times, asking which half holds the decisive mistake.
- **Root-cause attributor:** one call asking for a root cause, the earliest wrong step and a component.
- **RRSI trace digester** (Xia et al., 2026): RRSI’s own multi-turn failure digest; we take the earliest step it cites as evidence.
- **Counterfactual search** (CF search), replay-based like AgenTracer’s labelling (Zhang et al., 2026a) but without grading criteria: one call names up to five suspect steps with corrected actions; suspects are tested in order with 4 corrected and 4 null replays, and the search stops at the first correction that beats null by 2 successes or when its budget of 40 replays runs out; otherwise it answers with the first suspect.

Judges are prompted for the decisive mistake, defined as the earliest step whose correction would most likely have rescued the episode (Appendix B); post hoc reruns change this definition or withhold the grading information. Pro runs with thinking off unless stated.

Registration and analysis. Methods, prompts, metrics and analyses were registered, and all methods had been run, before the main profiles were computed; no method sees a profile. The registered primary comparison is counterfactual search at budget 40 against the all-at-once judge with Pro: the paired difference in $R(\hat{k})$ over rescuable failures, with a 95% bootstrap interval over task clusters (10,000 draws) and a paired permutation test flipping signs by task, on retail and airline pooled and, for AppWorld, separately. Everything else is estimated with cluster-bootstrap intervals and no further tests; unregistered analyses are marked post hoc or exploratory. A seeded 10% of failures in each domain got a second, independent profile. Costs are list-price dollars for tokens and replays. Appendix G lists every change made after registration.

5 Results

Profiles. Of 300 failures, 157 were rescuable: 65 of 100 in retail, 37 in airline and 55 in AppWorld. Null replays from the same step succeeded 27%, 18% and 11% of the time; without them, many rescues would be chance recoveries. A third of the rescuable failures (18, 11 and 26) had two or more steps with a gain of at least 0.5. In the retest, 27 of 30 failures got the same verdict on rescuability, 15 of the 17 that both runs found rescuable the same best rescue step (the other two one and two steps apart), and step-level gains correlated .63 to .80.

Under the registered prompts, the first-write rule scored highest. Table 1 gives every method’s rescue gain. Pooled over the three domains, the first write scored 0.42 [0.33, 0.51]; the LLM methods scored between 0.20 (step-by-step) and 0.34 (binary search). The first write’s lead over binary search was not significant (+0.08, exploratory, $[-0.02, +0.18]$); over the all-at-once judges it was +0.11 to +0.12. Exact agreement with the best rescue step ordered the methods the same way, and over all 300 failures, a cohort chosen without the reference, the first write scored 0.23 and the LLM methods 0.11 to 0.19 (Appendix Table 4). At the steps the methods named, the oracle proposed a correction about equally often (0.52–0.64 of samples); the first write led because corrections at its steps worked better (mean gain 0.61, against 0.34–0.51; post hoc, Appendix Table 9). Cost bought little: the all-at-once judges with Flash (\$0.001 per failure), Pro (\$0.003) and Pro with

thinking (\$0.014) scored 0.30, 0.30 and 0.31 (Appendix Figure 3). On Who&When our judges reproduce the published pattern (Appendix D).

Replays and grading information both helped search. Counterfactual search at budget 40 scored 0.26 pooled, at seven times the cost of the all-at-once judge with Pro (\$0.022 against \$0.003 per failure). The registered comparison with that judge detected no difference, with intervals that admit 0.1 either way: $-0.061 [-0.146, +0.032]$ in τ^2 (permutation $p = .20$, $n = 102$) and $+0.025 [-0.087, +0.136]$ in AppWorld ($p = .69$, $n = 55$). This comparison, and its cost ratio, belong to the registered configuration, in which search alone did not read the grading criteria. Post hoc, we reran search’s suspect call with the grading information the judges read, keeping its replays and budget (\$2.4), which completes a 2×2 of replays and information (Table 1, lower block). Search’s first suspect scored 0.22 without the information and 0.27 with it; after replays, 0.26 and 0.35. Information and replays raised the score by similar amounts (main effects $+0.07 [+0.01, +0.13]$ and $+0.06 [+0.03, +0.10]$), with no clear interaction ($+0.03 [-0.02, +0.09]$). Withholding it cost the judges about as much (0.07 and 0.11). At matched information search led the all-at-once judge by 0.04 $[-0.02, +0.09]$ without the grading and 0.05 $[-0.004, +0.11]$ with it, at five to seven times its cost. These contrasts are post hoc and descriptive.

Asked for the measured quantity, judges came close to the rule. Judges that missed the best rescue step usually named an earlier one (in retail, the all-at-once judge with Pro on 40 of 65 rescuable failures). The registered prompts asked for the earliest decisive mistake, while R_k measures the gain of a correction over letting the agent act again. Post hoc, we re-ran three judges asking instead for the step whose correction would most likely rescue the task, or would most increase success over letting the agent act again (Appendix B). The second, which matches the metric, raised the Pro judges by 0.07 (all-at-once 0.30 to 0.36, $[+0.02, +0.12]$; binary search 0.34 to 0.40, $[-0.01, +0.13]$; Table 1) and the Flash judge by 0.01. Binary search asked this way showed no detected difference from the first write ($-0.02 [-0.10, +0.07]$, an interval too wide to show equivalence), so the first write’s lead under the registered prompts owes much to how the judges were asked.

Method	Model	Mean rescue gain $R(\hat{k})$, rescuable failures [95%]				Exact (pooled)	\$ per failure
		Retail (65)	Airline (37)	AppWorld (55)	Pooled (157)		
First write	rule	0.53 [0.38, 0.65]	0.27 [0.13, 0.45]	0.39 [0.26, 0.52]	0.42 [0.33, 0.51]	0.38	0.0000
Last step	rule	0.11 [0.02, 0.24]	0.13 [0.04, 0.23]	0.36 [0.23, 0.50]	0.20 [0.13, 0.29]	0.13	0.0000
All-at-once	Flash	0.29 [0.19, 0.38]	0.41 [0.28, 0.53]	0.25 [0.14, 0.37]	0.30 [0.24, 0.37]	0.28	0.0010
All-at-once	Pro	0.30 [0.19, 0.41]	0.34 [0.21, 0.48]	0.26 [0.15, 0.38]	0.30 [0.23, 0.36]	0.31	0.0031
All-at-once	Pro, thinking	0.26 [0.17, 0.35]	0.42 [0.28, 0.57]	0.28 [0.18, 0.40]	0.31 [0.24, 0.38]	0.32	0.0141
Step-by-step	Pro	0.14 [0.03, 0.25]	0.16 [0.04, 0.25]	0.29 [0.18, 0.40]	0.20 [0.13, 0.26]	0.19	0.0099
Binary search	Pro	0.33 [0.21, 0.47]	0.34 [0.21, 0.46]	0.35 [0.24, 0.45]	0.34 [0.27, 0.41]	0.32	0.0064
Root-cause attributor	Pro	0.26 [0.16, 0.38]	0.36 [0.21, 0.53]	0.22 [0.10, 0.34]	0.27 [0.20, 0.34]	0.28	0.0031
RRSI trace digester	Pro	0.17 [0.07, 0.29]	0.35 [0.24, 0.49]	0.19 [0.11, 0.27]	0.22 [0.16, 0.29]	0.21	0.0077
Counterfactual search@40	Pro, replays	0.22 [0.12, 0.33]	0.31 [0.19, 0.44]	0.28 [0.17, 0.41]	0.26 [0.20, 0.34]	0.26	0.0222
<i>Post hoc variants</i>							
All-at-once, gain	Pro	0.39 [0.28, 0.51]	0.42 [0.28, 0.55]	0.29 [0.18, 0.42]	0.36 [0.29, 0.44]	0.36	0.0031
Binary search, gain	Pro	0.44 [0.33, 0.55]	0.42 [0.25, 0.60]	0.35 [0.23, 0.47]	0.40 [0.33, 0.48]	0.39	0.0071
All-at-once, no grading	Pro	0.17 [0.08, 0.26]	0.32 [0.19, 0.45]	0.25 [0.14, 0.35]	0.23 [0.17, 0.30]	0.24	0.0028
Binary search, no grading	Pro	0.21 [0.10, 0.33]	0.24 [0.12, 0.35]	0.24 [0.14, 0.35]	0.23 [0.16, 0.29]	0.20	0.0064
CF search@0	Pro	0.22 [0.14, 0.32]	0.13 [0.02, 0.24]	0.28 [0.17, 0.40]	0.22 [0.16, 0.28]	0.20	0.0037
CF search@0, grading	Pro	0.26 [0.17, 0.37]	0.24 [0.13, 0.34]	0.31 [0.18, 0.44]	0.27 [0.21, 0.34]	0.27	0.0037
CF search@40, grading	Pro, replays	0.34 [0.25, 0.45]	0.38 [0.26, 0.52]	0.34 [0.22, 0.46]	0.35 [0.28, 0.42]	0.33	0.0152

Table 1: Attribution scored by the rescue gain at the named step. Intervals: cluster bootstrap over tasks. Exact: share naming the best rescue step. Cost: mean over domains of list-price dollars per failure (tokens and replays). Lower block, post hoc (§5): judges asked for the largest gain over null (gain) or without grading information; search’s first suspect before any replay (@0) and search given the grading information.

The rule’s lead depends on the kind of step.

The first write led in retail (0.53, against 0.14–0.33 for LLM methods) and AppWorld (0.39); in airline the all-at-once judges led (0.41–0.42, against 0.27). Most best rescue steps were writes in retail (36 of 65), usually a wrong write replaced by the right one, and utterances in airline (25 of 37), typically a refusal or a clarifying question the agent should have made before acting. Split by the kind of best rescue step (post hoc, Appendix Table 14), the first write scored 0.60 on the 82 failures where that step changed state (a write, or AppWorld’s submission), against 0.46 for binary search asked for the largest gain (+0.13 [+0.03, +0.23]), and 0.23 on the other 75, against 0.34 to 0.39 for the best judges and informed search; 74% of its score came from the first group. Since the oracle repairs writes more readily than utterances (§7), the rule’s lead is conditional on this correction mechanism; with the eight domain-by-kind cells weighted equally, binary search asked for the largest gain scored 0.43 and the first write 0.41. Binary search was the only LLM method whose score held across domains (0.33–0.35).

The ranking held under most choices, but not without null replays. Appendix Table 6 re-scores every method under other choices (post hoc;

details in Appendix D). With fewer samples, thresholds from 0.25 to 0.75, only failures with a single rescuing step, or failures chosen on two samples and scored on the other two, the first write stayed first and binary search the best LLM method (Kendall’s τ of 0.78 to 1.00). With the observed action forced as the control, the judges gained most, and binary search asked for the largest gain scored 0.45 against the first write’s 0.44 (−0.01 [−0.10, +0.07]). A second oracle model (V4-Flash) found 120 failures rescuable instead of 157, agreed on the best rescue step for 71–72% of those both found rescuable, and still ranked the first write first ($\tau = .78$). A mechanical screen of the 4,795 corrections that entered a profile (Appendix Table 11) found an identifier, name or date the agent had not yet seen in 0.7% and two or more calls beyond the observed step in 1.8%; dropping them left the ranking unchanged ($\tau = 1.00$). Individual profiles are noisier than the ranking: halves of the data with two samples each agreed on the best rescue step for 56% of failures (Appendix F). Dropping the null replays reversed the ranking: 203 failures became rescuable, the first write fell to sixth place, and the all-at-once judge with thinking ranked first (0.42 against 0.37). Uncorrected replays from the first fifth of an episode succeeded 41–44% of the time and from the last fifth 2–6%

(Appendix Figure 2), and LLM methods other than binary search named steps on average 53–67% of the way through an episode, against 77% for best rescue steps, so without the null control they collect credit for recoveries that happen anyway.

Without the grading, the oracle rarely rescued writes. Choosing among plausible actions by a private criterion leaves no unseen value for a screen to find. Post hoc, we reran the oracle with the same instructions and context but no grading, telling it only that the episode failed, at the 960 steps any method named on the 157 rescuable failures (\$4.0; Appendix Table 15). Gains at best rescue steps fell from 0.76 to 0.17, and 28 of these steps kept a gain of 0.5. Corrections of writes lost nearly everything (retail 0.77 to 0.00, AppWorld 0.82 to 0.20), corrections of utterances half to two thirds (0.72–0.76 to 0.24–0.36): the oracle repairs writes mainly by knowing which write the task wanted. Under this oracle every method scored 0.03 to 0.11, the first write ranked seventh of the ten registered methods (0.05, against 0.11 for counterfactual search; $-0.06 [-0.10, -0.02]$), and the ordering kept little of the registered one ($\tau = .24$). The registered profile measures where a correction that knows the intended outcome helps, the information RRSI’s analyst has; for a critic without it, the first write’s lead does not hold.

6 Do Better Attributions Give Better Repairs?

Design. A rescue score is a proxy: a loop needs attributions only to repair its agent. We tested this on RRSI itself, in a separately registered experiment. We took two other runs of the loop in each τ^2 domain and the harness each run had accepted by the start of rounds 5 and 15, giving eight incumbent states. At each state we re-ran one unmodified RRSI round five times, changing only the evidence the proposer gets:

- **No attribution**: a note asking the proposer to read the traces itself.
- **RRSI analysis**: the round’s own analysis report and trace digests, unchanged.
- **First write, binary search (Pro) and counterfactual search@40**: for each failing trace, the step the method named and an excerpt of it, with no reason, component or method name (Appendix B).
- **Oracle step** (post hoc): the same digest naming the step with the largest rescue gain in the trace’s own profile, for the 60 traces where some step had a positive gain.
- **No reading** (post hoc): two groups whose proposer cannot open failing traces (a read returns a notice; passing traces stay readable), given either RRSI’s analysis, or for the same 60 traces the oracle step with its excerpt, the oracle’s diagnosis and the correction that rescued most often there (*oracle fix*).

The proposer works from the traces RRSI builds from the incumbent’s evaluation on training tasks (81 failing traces in all, 42 retail and 39 airline) and, as in RRSI, can read each of them, grading report included, except in the no-reading groups. Each round drafts two candidates and puts them through RRSI’s critic, smoke test and full evaluation on training tasks, after which RRSI’s own rule decides whether to accept one. A candidate is *deployable* if it passes these gates with a harness error rate of at most 2%. Every deployable candidate and every incumbent was deployed on 240 held-out episodes (retail 40 tasks $\times$ 6, airline 20 $\times$ 12).

Analysis. A candidate slot’s gain G is its held-out success minus its incumbent’s, in points, and 0 if it was not deployable, since the loop then keeps the incumbent. A contrast between two groups is the mean over the eight states of the difference in their mean gains. The six primary contrasts compare each step attribution with no attribution and with RRSI’s analysis; p values come from permuting group labels within states and are Holm-adjusted, and 95% intervals from a t distribution over states (df 7). A contrast whose 90% interval lies within ± 2.5 points, about one standard deviation of a single candidate’s gain in earlier runs, counts as practically equivalent. The margin reflects the noise of one deployment, not a threshold of practical value: over 20 rounds the four source runs moved held-out success by +7.9 points in one (airline, 90% [+2.5, +14.2]) and by -4.3 to $+1.7$ in the others, so a useful round gains well under a point. We therefore also report the smallest margin each interval excludes, and the outcome that matters to the loop, the gain of the candidate RRSI accepts. Grouping the eight states by their four source runs is a post hoc check. The design, with an expected standard error of 0.9 points per contrast, was registered before any attribution was run on these traces,

Evidence for the proposer	$R(\hat{k})$, these		Mean G , all 16 slots			Deployable		RRSI	Round
	retail / airline	Deployable	both	retail	airline	mean G	$G > 0$	accepted	level
No attribution	–	10/16	−4.38	−9.27	+0.52	−7.00	50%	2	−0.10
RRSI analysis	0.30 / 0.21	13/16	−0.34	−0.31	−0.36	−0.42	38%	1	−0.10
First write	0.37 / 0.16	10/16	−0.03	+0.21	−0.26	−0.04	60%	2	+0.47
Binary search	0.32 / 0.51	9/16	−0.05	+0.16	−0.26	−0.09	44%	0	0.00
Counterfactual search@40	0.27 / 0.32	12/16	−2.47	−0.26	−4.69	−3.30	25%	2	+0.16
Oracle step	0.67 / 0.81	10/16	−0.39	−1.04	+0.26	−0.62	50%	1	+0.26
RRSI analysis, no reading	0.30 / 0.21	14/16	−0.39	−0.42	−0.36	−0.45	21%	2	−0.05
Oracle fix, no reading	0.67 / 0.81	10/16	0.00	+0.42	−0.42	0.00	60%	1	+0.10

Primary contrast	Difference [95%]	p	Holm	Equivalent	Round level [95%]
First write – none	+4.35 [−5.97, +14.67]	0.52	1.00	no	+0.57 [−0.19, +1.34]
Binary search – none	+4.32 [−6.90, +15.55]	0.66	1.00	no	+0.10 [−0.14, +0.35]
CF search@40 – none	+1.90 [−10.24, +14.04]	0.91	1.00	no	+0.26 [−0.99, +1.51]
First write – RRSI	+0.31 [−0.56, +1.18]	0.65	1.00	yes	+0.57 [−0.32, +1.46]
Binary search – RRSI	+0.29 [−0.82, +1.39]	0.60	1.00	yes	+0.10 [−0.14, +0.35]
CF search@40 – RRSI	−2.14 [−5.67, +1.40]	0.31	1.00	no	+0.26 [−0.99, +1.51]
Oracle – RRSI	−0.05 [−1.97, +1.87]	0.98	–	yes	+0.36 [−0.27, +0.99]
Oracle – step groups	+0.46 [−1.62, +2.54]	0.75	–	yes	–
Oracle fix – RRSI, no reading	+0.39 [−0.98, +1.76]	0.53	–	yes	+0.16 [−1.44, +1.76]
RRSI: no reading – reading	−0.05 [−1.31, +1.21]	0.97	–	yes	+0.05 [−1.55, +1.66]
Oracle fix, no reading – oracle step	+0.39 [−0.78, +1.56]	0.61	–	yes	−0.16 [−0.85, +0.54]

Table 2: Repairs drafted from each kind of evidence in τ^2 : 8 states $\times$ 2 candidates per group. G : held-out success of a candidate minus its incumbent’s, in points, 0 if not deployable. $R(\hat{k})$: rescue gain of the steps a group pointed to, on the 33 rescuable failures among the experiment’s 81 failing traces (post hoc; for RRSI, the earliest step its digests cite). Accepted: by RRSI’s own rule. Round level: gain of the accepted candidate, 0 if none (tested by exact sign flips, text). Contrasts: mean over states of group differences, p from within-state permutation; Equivalent: 90% interval within ± 2.5 points. The oracle and no-reading groups and their contrasts are post hoc and not Holm-adjusted.

and no held-out result was looked at until all 62 deployments had finished.

No contrast was significant. Of 80 drafts, 54 were deployable, 9 to 13 per group (Table 2); none failed the error check. No primary contrast was significant (Holm-adjusted $p = 1.0$ for all six). Repairs drafted from first-write and binary-search steps were practically equivalent to those drafted from RRSI’s analysis: +0.31 [−0.56, +1.18] and +0.29 [−0.82, +1.39] points; their 90% intervals exclude differences beyond 1.0 and 1.2 points, or 1.5 and 1.4 points with the states grouped by run, but not differences the size of a typical round’s gain. Counterfactual search was 2.14 points below RRSI’s analysis [−5.67, +1.40], mostly through two airline candidates that lost 23 and 14 points. The contrasts with no attribution are dominated by one candidate drafted without attribution whose write gate blocked nearly every write (−72.5 points; without it the group’s mean gain was +0.17).

The step groups differed little in rescue score.

Post hoc, we computed rescue profiles for the 81 failing traces these proposers saw (\$5.0); only 33 were rescuable (17 retail, 16 airline). On them the first write scored 0.27, binary search 0.41, counterfactual search 0.30 and the earliest step RRSI’s own digests cite 0.26 (Table 2): the first write was no better a pointer here than RRSI’s analysis (+0.01 [−0.15, +0.19]), and binary search better but not significantly (+0.15 [−0.03, +0.35]). The rescue score of the pointers a proposer received correlated only weakly with its candidates’ gains (Spearman .15 within states, $p = .14$). In no group was the mean gain above zero (21–60% of deployable candidates gained). Rounds rarely accepted a candidate (0 to 2 of 8 per group, 11 of 64 in all), and RRSI’s own rule rejected all three that lost more than 10 points, so the gain of the accepted round was usually 0 and differed little between groups (post hoc with states grouped by run, Appendix Table 13): against RRSI’s analysis, 0.00 to +0.57 points, all 95% intervals within ± 2.7 ; rounds given any step pointer, +0.33 [−0.39, +1.04].

An oracle pointer did not help either. To test a pointer that scores highest by construction, we ran a sixth group post hoc with the step of largest rescue gain (0.67 and 0.81 on rescuable failures in retail and airline; \$5.9), keeping its held-out results unseen until its 10 deployments had finished. Its candidates’ mean gain was -0.39 points, practically equivalent to RRSI’s analysis ($-0.05 [-1.97, +1.87]$; within ± 2.1 with states grouped by run) and to the three registered step groups pooled ($+0.46 [-1.62, +2.54]$, though not once grouped by run). Its pointers were chosen on the replays that scored them; with fresh oracle samples and replays (post hoc, \$2.3; Appendix Table 16) their gain on rescuable failures fell from 0.74 to 0.59, still above binary search’s 0.42 and RRSI’s 0.16.

Without reading, an oracle’s diagnosis and correction did not help either. To give proposers evidence they could not find themselves, we ran the two no-reading groups post hoc (\$12.4), keeping their held-out results unseen until all their deployments had finished. Blocking the reads did not change the repairs drafted from RRSI’s analysis ($-0.05 [-1.31, +1.21]$ points against the same evidence with reading), and the oracle’s step, diagnosis and correction did not improve on RRSI’s analysis ($+0.39 [-0.98, +1.76]$; 90% interval within ± 1.5 , or ± 2.0 with states grouped by run), although on fresh replays the corrections it was given kept a gain of 0.65 of their 0.85 (at least 0.5 for 70% of the 60 traces).

Proposers took up the evidence but did not fix the cases it named. The proposers’ logs and drafts show where the evidence was lost (post hoc). In every group that could read them, the proposer opened 97–100% of the failing traces, and without any attribution 95–97% of the reads of a failing trace covered the step each pointer named. Each draft lists the training tasks it expects to fix; given the oracle’s evidence, drafts named 35–42% of the traces that carried it, against 22–28% from RRSI’s analysis ($+16$ points over states, $[-6, +39]$). But in RRSI’s evaluation of the drafts on the training tasks, the failing tasks a draft named improved no more than those it did not name (-3.3 points $[-10.0, +3.5]$ over 87 drafts; $-12.5 [-26.0, +1.0]$ with the oracle’s evidence), and a draft’s gain on failing training tasks tracked its held-out gain only weakly (Spearman .19, $p = .08$). The evidence was lost in turning a diagnosis into an edit that

fixes even the case it targets, before any question of transfer.

7 Discussion

What the score measures. R_k belongs to a step and an oracle: how often the oracle proposes a correction, times how well its corrections work. A low R_k can mean that this oracle could not repair a step, and a high R_k marks an effective point of recovery for an oracle that knows the intended outcome, which need not be where the failure began; with the grading the oracle repaired writes best (in retail 0.52 against 0.13 for utterances; Appendix Table 10), without it almost never. The profile is a reference for intervention, not a record of where the error lies, and rankings built on it depend on the null control and on what the oracle may know; replay-based benchmarks should report both. Profiles ranked methods consistently, but individual best rescue steps are noisy (split halves agreed on 56%), so reusing them as per-failure labels calls for gains estimated on held-out samples.

Implications for self-improving loops. In this loop the evidence interventions brought no detectable incremental benefit, and the drafts suggest why: proposers took up even the oracle’s step, diagnosis and correction, but their edits did not fix the cases it named (§6), and most edits in every group moved held-out success by less than 5 points (Appendix Figure 4). With no positive control showing that this proposer turns useful evidence into gains at these states, we cannot separate weak attribution utility from an insensitive repair process; a bottleneck after attribution remains a hypothesis. RRSI’s re-evaluation rejected every edit that misfired broadly: the evaluation that decides which repairs to keep protected the agent more than the attribution behind them.

8 Conclusion

Against counterfactual rescue with null controls and an oracle that reads the grading, a free structural rule was competitive with LLM attribution; its lead came from writes the oracle could repair only with the grading, and it disappeared for an oracle without it. Baselines, the definition of the target and the information given to methods and oracle moved the comparison as much as the method did. Neither the step of largest rescue gain nor its diagnosis and correction, which kept most of their gain on fresh replays, detectably improved the repairs of

the loop we studied. We release failures, profiles, replay outcomes and code.

Limitations

All agents, oracles and attributors use one model family (DeepSeek-V4); the second oracle is a smaller model of the same family. The failures come from one self-improvement loop and three domains, with 37 to 65 rescuable failures each, so per-domain intervals are wide. The reference rests on one kind of intervention, a single action corrected by an oracle that sees the grading: it says where a correction informed by the intended outcome helps, not where an error began. Without the grading the oracle rarely rescued writes and the ranking changed; we ran that oracle only at steps some method named, so it gives no full profile. A mechanical screen found few corrections that use unseen values or do more than one step, but it cannot detect a choice made with the grading, and no human has audited the corrections; we release a stratified sample of 60 with an audit sheet. Individual best rescue steps are noisy: split halves agreed on 56%. The threshold of 0.5 and four samples per step were fixed on a 20-failure pilot. Many analyses are post hoc: the robustness checks, the observed-action control, the second oracle, the oracle without grading, the split by kind of step, the admissibility screen, the prompt variants, the judges without grading information, the informed search, the gold-deviation rule, the reading logs, and in the repair experiment the rescue scores of its pointers and their fresh replays, its run-grouped and accepted-round intervals, and its oracle and no-reading groups. Our registered prompts asked for the earliest decisive mistake while the metric rewards the largest gain over null, so the registered comparison understates the judges. Our counterfactual search labels steps by replay but trains no attributor, unlike AgenTracer (Zhang et al., 2026a), so its results do not speak to trained attributors. Fourteen retail failures were drawn after the others, from the two further runs of the loop that the repair experiment also used (one trace is in both). Two of 6,540 AppWorld replays could not run (Appendix G). Who&When is a multi-agent benchmark with a different format; only methods without replay were run there. The repair experiment covers τ^2 only, with eight states from four runs of the loop, 16 candidate slots per group and 33 rescuable failures among its 81 failing traces. Its

margin of 2.5 points reflects the noise of one deployment and exceeds a typical round’s gain, and its intervals widen when states are grouped by run; it rules out large differences from RRSI’s analysis, not small ones that would compound over a long loop. Only 11 of 64 rounds accepted a candidate, and its contrasts with no attribution are dominated by one candidate. The registered digests carried only the step; the post hoc groups added a diagnosis and a correction, but no group named a harness component to change, and all used one loop and one proposer model.

Ethics Statement

All environments use synthetic users and data. Attribution methods that misdiagnose failures can lead self-improving agents to change the wrong behaviour; our results suggest that such diagnoses should be checked against simple baselines before they drive automatic changes.

References

- Victor Barres, Honghua Dong, Soham Ray, Xujie Si, and Karthik Narasimhan. 2025. τ^2 -bench: Evaluating conversational agents in a dual-control environment. *arXiv preprint arXiv:2506.07982*.
- Akash Bonagiri, Devang Borkar, Gerard Janno Andrias, Setareh Rafatirad, and Houman Homayoun. 2026. CausalFlow: Causal attribution and counterfactual repair for LLM agent failures. *arXiv preprint arXiv:2605.25338*.
- Mert Cemri, Andrei Cojocaru, Melissa Pan, Shu Liu, Shubham Agarwal, Alexander Krentsel, Jay Tang, Kannan Ramchandran, Joseph E. Gonzalez, Matei Zaharia, Alexandros G. Dimakis, and Ion Stoica. 2026. Fantastic adaptive taxonomies and how to use them. *arXiv preprint arXiv:2607.16387*.
- Mert Cemri, Melissa Z. Pan, Shuyi Yang, Lakshya A. Agrawal, Bhavya Chopra, Rishabh Tiwari, Kurt Keutzer, Aditya Parameswaran, Dan Klein, Kannan Ramchandran, Matei Zaharia, Joseph E. Gonzalez, and Ion Stoica. 2025. Why do multi-agent LLM systems fail? *arXiv preprint arXiv:2503.13657*.
- Mengzhuo Chen, Junjie Wang, Zhe Liu, Yawen Wang, Haiming Zheng, and Qing Wang. 2026. From failed trajectories to reliable LLM agents: Diagnosing and repairing harness flaws. *arXiv preprint arXiv:2606.06324*.
- Joseph Y. Halpern. 2016. *Actual Causality*. MIT Press.
- Yeonjun In et al. 2026. Rethinking failure attribution in multi-agent systems: A multi-perspective benchmark and evaluation. In *International Conference on Machine Learning*. ArXiv:2603.25001.

776	Jiazheng Li, Emine Yilmaz, Bei Chen, and Dieu-	2026b. Agentic context engineering: Evolving con-	831
777	Thu Le. 2026. Towards self-improving error di-	texts for self-improving language models. In <i>Inter-</i>	832
778	agnosis in multi-agent systems. <i>arXiv preprint</i>	<i>national Conference on Learning Representations</i> .	833
779	<i>arXiv:2604.17658</i> .		
780	Xiaofeng Lin, Yingxu Wang, Tung Sum Thomas Kwok,	Shaokun Zhang, Ming Yin, Jieyu Zhang, Jiale Liu,	834
781	Daniel Guo, Sahil Arun Nale, Charles Fleming,	Zhiguang Han, Jingyang Zhang, Beibin Li, Chi	835
782	and Guang Cheng. 2026. REFLECT: Intervention-	Wang, Huazheng Wang, Yiran Chen, and Qingyun	836
783	supported error attribution for silent failures in LLM	Wu. 2025. Which agent causes task failures and	837
784	agent traces. <i>arXiv preprint arXiv:2606.09071</i> .	when? on automated failure attribution of LLM	838
785	Jiale Liu, Huajun Xi, Shaokun Zhang, Yifan Zeng, Tian-	multi-agent systems. In <i>Proceedings of the 42nd</i>	839
786	wei Yue, Chi Wang, Jian Kang, Qingyun Wu, and	<i>International Conference on Machine Learning</i> .	840
787	Huazheng Wang. 2026. Who&when pro: Can LLMs	ArXiv:2505.00212.	841
788	really attribute failures in AI agents? <i>arXiv preprint</i>	Andrew Zhao, Daniel Huang, Quentin Xu, Matthieu	842
789	<i>arXiv:2607.09996</i> .	Lin, Yong-Jin Liu, and Gao Huang. 2024. ExpeL:	843
790	Ming Ma, Jue Zhang, Fangkai Yang, Yu Kang, Qing-	LLM agents are experiential learners. In <i>Proceedings</i>	844
791	wei Lin, Saravan Rajmohan, and Dongmei Zhang.	<i>of the AAAI Conference on Artificial Intelligence</i> .	845
792	2025. DoVer: Intervention-driven auto debug-	Kunlun Zhu, Zijia Liu, Bingxuan Li, Muxin Tian, Yingx-	846
793	ging for LLM multi-agent systems. <i>arXiv preprint</i>	uan Yang, Jiaxun Zhang, Pengrui Han, Qipeng Xie,	847
794	<i>arXiv:2512.06749</i> .	Fuyang Cui, Weijia Zhang, Xiaoteng Ma, Xiaodong	848
795	Judea Pearl. 2009. <i>Causality: Models, Reasoning, and</i>	Yu, Gowtham Ramesh, Jialian Wu, Zicheng Liu, Pan	849
796	<i>Inference</i> , 2nd edition. Cambridge University Press.	Lu, James Zou, and Jiaxuan You. 2025. Where LLM	850
797	Jaineet Shah. 2026. Causal agent replay: Counterfactual	agents fail and how they can learn from failures.	851
798	attribution for LLM-agent failures. <i>arXiv preprint</i>	<i>arXiv preprint arXiv:2509.25370</i> .	852
799	<i>arXiv:2606.08275</i> .	Kunlun Zhu, Xuyan Ye, Zhiguang Han, Yuchen Zhao,	853
800	Noah Shinn, Federico Cassano, Ashwin Gopinath,	Bingxuan Li, Weijia Zhang, Muxin Tian, Xiangru	854
801	Karthik Narasimhan, and Shunyu Yao. 2023. Re-	Tang, Pan Lu, James Zou, Jiaxuan You, and Heng Ji.	855
802	flexion: Language agents with verbal reinforcement	2026. AgentDebugX: An open-source toolkit for fail-	856
803	learning. In <i>Advances in Neural Information Pro-</i>	ure observability, attribution, and recovery in LLM	857
804	<i>cessing Systems</i> .	agents. <i>arXiv preprint arXiv:2607.18754</i> .	858
805	Harsh Trivedi, Tushar Khot, Mareike Hartmann, Ruskin	A Profile Details	859
806	Manku, Vinty Dong, Edward Li, Shashank Gupta,	Replays. τ^2 replays restore the database state	860
807	Ashish Sabharwal, and Niranjan Balasubramanian.	after the kept steps and continue with the same	861
808	2024. AppWorld: A controllable world of apps and	agent harness and user simulator (DeepSeek-V4-	862
809	people for benchmarking interactive coding agents.	Flash, temperature 0); a corrected action replaces	863
810	In <i>Proceedings of the 62nd Annual Meeting of the</i>	the agent’s turn at step k (a message, or tool calls	864
811	<i>Association for Computational Linguistics</i> .	whose results are then computed). AppWorld re-	865
812	Rakshith Vasudev, Melisa Russak, Dan Bikel, and	plays restore the app state by re-executing the kept	866
813	Waseem Alshikh. 2026. Accurate failure prediction	code and force the corrected turn, which must con-	867
814	in agents does not imply effective failure prevention.	tain one Python block; the episode continues with	868
815	<i>arXiv preprint arXiv:2602.03338</i> .	the incumbent harness. A correction the replay	869
816	Peng Xia, Rujun Han, Zifeng Wang, Yanfei Chen, Yu-	cannot apply counts as giving no gain.	870
817	fan Zhuang, Yoonho Lee, Chengsong Huang, Han	From earliest step to profile. Our first registered	871
818	Yu, Zhongying CuiZhu, Yifei Ming, Huaxiu Yao, Bu-	protocol named the earliest step whose correction	872
819	rak Gokturk, Tomas Pfister, and Chen-Yu Lee. 2026.	succeeded on at least 2 of 4 replays more than 4 null	873
820	RRSI: Regularized recursive self-improvement of	replays did. On 20 pilot failures, independent runs	874
821	agent harnesses. <i>arXiv preprint arXiv:2609.24972</i> .	of it named the same step in 42% of pairs, mostly	875
822	Guibin Zhang, Junhao Wang, Junjie Chen, Wangchun-	because the oracle corrected a step in one run and	876
823	shu Zhou, Kun Wang, and Shuicheng Yan. 2026a.	not in another. Before computing any profile on	877
824	AgentTracer: Who is inducing failure in the LLM	the main set, we registered protocol 3 instead.	878
825	agentic systems? In <i>International Conference on</i>	Retest. The retest failures were drawn with a	879
826	<i>Learning Representations</i> . ArXiv:2509.03312.	seeded random choice of 10% per domain before	880
827	Qizheng Zhang, Changran Hu, Shubhangi Upasani,	any profile was computed and run with fresh oracle	881
828	Boyuan Ma, Fenglu Hong, Vamsidhar Kamanuru,	samples and replays.	882
829	Jay Rainton, Chen Wu, Mengmeng Ji, Hanchen Li,		
830	Urmish Thakker, James Zou, and Kunle Olukotun.		




Figure 1: Rescue profile of one retail failure: mean success of the oracle’s corrections and of uncorrected (null) replays, run only where some sample corrected. Steps 8–10 rescue; D: best rescue step, F: first write, J: all-at-once judges, B: binary search. At step 6, corrected and null replays both succeeded three times in four.




Figure 2: Success of uncorrected (null) replays from a step, by the step’s position in the failed episode (fifths), at the steps where the oracle proposed a correction. Each point averages 27 to 35 steps in the first fifth and 81 to 323 in the others.

Cost. The rescue profiles cost \$11.1 for the first 186 τ^2 failures and \$13.7 for AppWorld (plus \$0.7 for its retest); all parts of the 14 τ^2 failures added later cost about \$1.5. The attribution study, including methods, search, pilots and Who&When, cost about \$58, and the post hoc analyses for the main set \$23.1: judges with other definitions or without grading information \$9.4, the observed-action control \$3.9 and the second oracle \$9.8. The repair experiment (§6) cost about \$35: \$6.4 for drafting and critique, \$1.8 for attribution and search replays, \$7.9 for RRSI’s evaluations and smoke tests (5,639 episodes) and \$18.9 for held-out deployment (14,880 episodes); post hoc, its failures’ rescue profiles cost \$5.0 and the oracle-pointer group \$5.9.

B Prompts

The τ^2 prompts follow, verbatim; braces mark what is filled in. The AppWorld prompts replace the description of the agent and the task, the grading report and the component descriptions with AppWorld’s (the code agent, its unit tests, and the same nine components worded for it), and the oracle and search ask for one assistant turn with one Python

block.

Shared introduction (all judges)

A customer-service agent (an LLM inside a harness) failed a task of tau2-bench. Below is the failure as the developers’ analysis tools see it: the customer’s hidden instructions, the conversation (each agent step is numbered), and the grading.

The DECISIVE MISTAKE is the earliest agent step such that, had the agent acted correctly at that step and continued normally, the task would most likely have succeeded.

Harness components (a fix to the harness changes one of these):

- prompt: the system prompt text (instructions and policy)
- wording the model reads)
- control_flow: the agent loop code: how the next action is produced, checks or gates applied to tool calls before they run, retries, enforced confirmations
- config: numeric settings such as temperature, token limits or step limits
- output_plumbing: how tool results and tool descriptions are rendered to the model
- context_mgmt: which parts of the conversation history the model sees
- client_tool: extra tools the harness gives the model (e.g. a checker or calculator)
- skill: instruction files loaded for specific situations
- memory: structured state the harness keeps across steps
- subagent: an extra model call, e.g. a second model reviewing an action

```
=== FAILURE ===
{the failure: hidden instructions, numbered steps, grading
  report}
```

All-at-once

Find the decisive mistake and the harness component whose change would most likely have prevented it. Answer with JSON:

```
{
  "step": <step number>,
  "top3": [<the three most likely step numbers, best first>],
  "component": "<component>",
  "components3": [<three most likely components, best first>],
  "reason": "<one sentence>"
}
```

Step-by-step (asked for $k = 0, 1, \dots$ on the episode up to step k)

Judge step {k} only, given what happened before it. Is the agent’s action at step {k} the decisive mistake (the earliest step whose correction would most likely have made the task succeed)? Answer with JSON:

```
{
  "decisive": true | false,
  "component": "<component, if decisive>",
  "reason": "<one sentence>"
}
```

Binary search (episode up to step hi)

The decisive mistake is assumed to lie between step {lo} and step {hi}. Is it in the first half (steps {lo} to {mid}) or in the second half (steps {mid + 1} to {hi})? Answer with JSON:

```
{
  "half": "first" | "second",
  "component": "<the component most likely at fault>",
  "reason": "<one sentence>"
}
```

Root-cause attributor

An LLM agent failed a task. Find the root cause of the failure.

Agent components (id: role):

```
{COMP_TEXT}
```

Trace:

```
{view(rec)}
```

Give (1) a one-sentence root cause describing the cause, not the symptom;

(2) the index of the earliest step where the agent went wrong;

(3) the single component whose change would most plausibly prevent this failure.

Answer with JSON: {"root_cause": "...", "step": <int>, "component": "<component id>"}

Oracle (system prompt) As run; its opening sentence predates our terminology, and the profile it builds is the oracle-dependent reference of §3, not a ground truth about where a failure began.

You construct ground truth for failure attribution. A
→ customer-service agent failed a task of tau2-bench. You see
→ the domain policy, the agent's tools, the task's grading
→ criteria (the database must end as if exactly the gold
→ actions were executed, and the agent must tell the user
→ every required piece of information), and the conversation
→ up to one step of the agent.

Decide whether the agent's action at that step is a mistake: an
→ action a careful agent following the policy would not take
→ on the way to the gold outcome. Actions that are correct,
→ or harmless and do not lead away from the gold outcome, are
→ ok.

If it is a mistake, give the single action a careful agent would
→ take at that step instead: one assistant turn, either a
→ message to the user or one or more tool calls. Correct only
→ this step; do not do more than a careful agent would do at
→ this point. The corrected action must be one the agent
→ could have taken with what it knew then (the conversation
→ and tool results so far, and the policy); use the grading
→ criteria only to know what the right outcome is, never to
→ supply ids, amounts or facts the agent had not yet learned.

Answer with JSON:
{ "verdict": "ok" | "mistake", "why": "<one sentence>",
 "action": { "message": "<text to the user>" | { "tool_calls":
 → [{"name": "<tool>", "arguments": {...}]] } }
 ("action" only when the verdict is mistake.)

Counterfactual search, suspect call (system prompt)

A customer-service agent on tau2-bench failed its task: the
→ customer's request was not resolved as the domain policy
→ requires. You see the domain policy, the agent's tools and
→ the conversation; you do not know what the correct outcome
→ was.

Name the steps where the agent most likely went wrong, most
→ suspect first (at most 5). For each, give the single action
→ a careful agent would have taken at that step instead (one
→ assistant turn: a message to the user or one or more tool
→ calls, using only what the agent knew at that point) and
→ the harness component most likely at fault.

Harness components:
- prompt: the system prompt text (instructions and policy
→ wording the model reads)
- control_flow: the agent loop code: how the next action is
→ produced, checks or gates applied to tool calls before they
→ run, retries, enforced confirmations
- config: numeric settings such as temperature, token limits or
→ step limits
- output_plumbing: how tool results and tool descriptions are
→ rendered to the model
- context_mgmt: which parts of the conversation history the
→ model sees
- client_tool: extra tools the harness gives the model (e.g. a
→ checker or calculator)
- skill: instruction files loaded for specific situations
- memory: structured state the harness keeps across steps
- subagent: an extra model call, e.g. a second model reviewing
→ an action

Answer with JSON:
{ "suspects": [{ "step": <int>, "why": "<one sentence>",
 → "component": "<component>",
 "action": { "message": "<text>" | { "tool_calls": [{ "name":
 → "<tool>", "arguments": {...}]] } } }

Variants (post hoc). The reruns of the all-at-once judges and binary search in §5 use the same prompts with one sentence of the introduction replaced. Registered:

The DECISIVE MISTAKE is the earliest agent step such that, had
→ the agent acted correctly at that step and continued
→ normally, the task would most likely have succeeded.

Most likely rescue:

The DECISIVE MISTAKE is the agent step whose correction would
→ most likely have rescued the task: had the agent acted
→ correctly at that step and continued normally, the task
→ would most likely have succeeded. If several steps qualify,
→ choose the one where a correction is most likely to make
→ the task succeed, which need not be the earliest.

Largest gain over null:

The DECISIVE MISTAKE is the agent step where a correction would
→ make the biggest difference: the step at which acting
→ correctly, and then continuing normally, would most
→ increase the chance that the task succeeds compared with
→ letting the agent act again from that step without any
→ correction. A step from which the task would often have
→ succeeded anyway, without a correction, is a weak choice
→ even if a correction there would succeed.

The no-grading variant keeps the registered definition, removes the grading report and the customer's hidden instructions from the episode, gives the domain policy and tools (AppWorld: the agent's rules and the APIs) as counterfactual search receives them, and opens with: "... You see the domain policy and the agent's tools, but not the customer's instructions or the grading, so you do not know what the correct outcome was."

Repair experiment: notes to the proposer. RRSI's proposer prompt is unchanged. In the no-attribution group the round's analysis report is replaced by the first note below; the step groups get the second note and one digest per failing trace, { "task_id", "lens": "failure", "decisive_step": *k*, "step_excerpt": <the step, at most 600 characters>, "blocker": "failure decided at step *k*" }.

No analysis report this round. Read the traces yourself
→ (list_traces, then read_trace) to find what went wrong; the
→ failing traces are listed first.

No analysis report this round. Read the traces yourself
→ (list_traces, then read_trace) to find what went wrong; the
→ failing traces are listed first. For each failing trace,
→ PER-TASK DIGESTS give decisive_step: the step at which an
→ automatic failure-attribution method located the failure.
→ Read the trace around that step (read_trace with
→ from_step/to_step) to see what went wrong there.

C Secondary Metrics, Step Kinds and Search Budgets

Table 4 gives the secondary metrics, Table 3 the kinds of agent steps and best rescue steps, and Table 5 counterfactual search by replay budget. Search's suspects included a step with a gain of at least 0.5 on 122 of 157 rescuable failures (78%); its confirmations were the weak point. Of the 87 rescuable failures where its own test (4 corrected against 4 null replays) confirmed a suspect, the

τ^2	steps/dec.	utterance	read	write	transfer
Retail	1165/65	54%/28	34%/1	11%/36	1%/0
Airline	1026/37	55%/25	31%/4	12%/8	1%/0
AppWorld	steps/dec.	docs	read	write	submit
AppWorld	1961/55	27%/3	45%/9	14%/29	5%/9

Table 3: Share of all agent steps of each kind / number of best rescue steps of that kind. AppWorld "docs" looks up API documentation; 5 AppWorld best rescue steps (9% of steps) are other code.




Figure 3: Rescue score against cost. Circles: all-at-once judges; squares: other LLM methods; triangles: counterfactual search at budgets 8, 16, 24 and 40. The free first-write rule lies above every method.

confirmed step had a gain of at least 0.5 in 41; elsewhere its corrections at the right step failed, or a weaker step passed the test first.

D Post-hoc Variants and Who&When

Robustness details. Choosing rescuable failures and scoring on the same noisy gains flattered the best rescue step (0.75, against 0.48 out of sample) but no method over another. With the observed action forced as the control, scored on the fixed cohort of 157 registered rescuable failures (156 of which would also be rescuable under this control), the all-at-once judge with Pro rose from 0.30 to 0.39, leaving the first write ahead of the best registered judge by 0.04 $[-0.07, +0.15]$. The second oracle proposed corrections less often (0.16–0.26 of samples against 0.22–0.37), and the two oracles agreed on rescuability for 69–80% of failures ($\kappa = .36$ – $.61$). Under its profile the first write led binary search by $+0.08$ $[+0.01, +0.16]$, and under the average of both profiles by $+0.10$ $[+0.02, +0.19]$ ($\tau = .82$); binary search asked for the largest gain came within 0.02 of it under both. The admissibility screen also flagged a number not seen before in 6.1% of corrections, mostly sums or differences the agent could compute; corrections of retail writes, the most effective kind, were flagged for an un-

seen value in none and for any unseen value in 0.4%. Without the corrections flagged for an unseen value or two more calls, 152 failures were rescuable, the first write scored 0.43 and binary search 0.34 ($+0.09$ $[-0.01, +0.19]$), and binary search asked for the largest gain came within 0.02 ($+0.02$ $[-0.07, +0.10]$). Without null replays a step scores $R_k + (m_k/K) \hat{p}_k^{\text{null}}$, where m_k of the K samples corrected it; at the registered rescuable failures the added term averaged 0.08 to 0.13 for the LLM methods and 0.04 for the first write.

Table 7 compares the judges’ registered prompts with the re-prompted variants (Appendix B), Table 8 a rule naming the first deviation from τ^2 ’s gold actions, which the judges read (0.50, though the oracle corrects from the same gold actions), Tables 9 and 10 split rescue gains into how often the oracle proposed a correction and how well its corrections worked, and Table 11 screens the corrections for information the agent had not seen and for scope. A stratified sample of 60 corrections (20 per domain, half of them flagged, spread over kinds of step) is released with an audit sheet asking whether each uses unavailable information, does more than one careful step, and is a step a careful agent could take.

Who&When. To check that our judges are not unusually weak, we ran the all-at-once, step-by-step and binary-search judges on the 184 logs of Who&When with its authors’ prompts and the setting that shows the judge the task’s correct final answer (Who&When’s “ground truth”, not an attribution label) (Zhang et al., 2025). They identified the responsible agent in 58–67% of logs with all-at-once judging and the responsible step in 3–25% (Table 12), the same pattern the benchmark reports: agents are found far more often than steps.

E Repair Candidates

Figure 4 shows the held-out gain of every deployable candidate in the repair experiment.

F Further Checks

All analyses and runs in this section are post hoc.

Admissibility by kind of step. At the first write’s steps on rescuable failures, no correction used an identifier, name or date the agent had not seen, 0 to 4.5% used an unseen number (airline writes the most), and 0 to 3% added two or more calls. At best rescue steps one of 553 corrections used an unseen

Method	Model	Exact [95%]	Within 1	Earliest	Share of max	$R(\hat{k})$, all failures
First write	rule	0.38 [0.27, 0.49]	0.62	0.33	0.52	0.23 [0.17, 0.30]
Last step	rule	0.13 [0.07, 0.20]	0.43	0.08	0.25	0.11 [0.07, 0.16]
All-at-once	Flash	0.28 [0.20, 0.36]	0.52	0.31	0.40	0.19 [0.14, 0.23]
All-at-once	Pro	0.31 [0.22, 0.39]	0.52	0.32	0.39	0.17 [0.13, 0.22]
All-at-once	Pro, thinking	0.32 [0.22, 0.42]	0.60	0.36	0.42	0.18 [0.14, 0.23]
Step-by-step	Pro	0.19 [0.13, 0.26]	0.38	0.22	0.25	0.11 [0.07, 0.15]
Binary search	Pro	0.32 [0.24, 0.40]	0.58	0.34	0.43	0.19 [0.15, 0.24]
Root-cause attributor	Pro	0.28 [0.19, 0.37]	0.49	0.31	0.36	0.16 [0.12, 0.20]
RRSI trace digester	Pro	0.21 [0.14, 0.29]	0.41	0.22	0.29	0.14 [0.10, 0.17]
Counterfactual search@40	Pro, replays	0.26 [0.18, 0.35]	0.44	0.31	0.35	0.15 [0.11, 0.19]

Table 4: Secondary metrics, pooled over the three domains. Exact, within one step, earliest rescuing step and share of the largest gain are over rescuable failures; the last column is over all 300 failures.

Budget	Retail	Airline	AppWorld	Pooled	\$
8	0.22	0.13	0.28	0.22	0.0125
16	0.21	0.23	0.27	0.23	0.0184
24	0.22	0.31	0.27	0.26	0.0212
32	0.22	0.31	0.28	0.26	0.0222
40	0.22	0.31	0.28	0.26	0.0222

Table 5: Counterfactual search: mean $R(\hat{k})$ on rescuable failures by replay budget, read off the same runs, and dollars per failure.




Figure 4: Held-out gain of every deployable candidate in the repair experiment, by the evidence its proposer got (*: no reading of failing traces; Fix: oracle step, diagnosis and correction). Filled: accepted by RRSI’s own rule. One candidate (−72.5) is drawn at −25.

identifier (an AppWorld write); two or more extra calls were most common at airline utterances (12% of 85), where a correction replaced a message by several lookups, and absent at retail steps.

Intervention scope in AppWorld. A correction replaces one code block by another of similar size: 5.4 lines against 5.3 in the observed block (median 3, 90th percentile 13) and 1.04 API calls against 0.96; 97% have at most one call more than the observed block (or than one call). At best rescue steps 78% of a correction’s lines are new. A corrected block thus often does more than a conversational

Reference	Resc.	First wr.	Best LLM	Rank	τ
Registered	157	0.42	Binary search 0.34	1	1.00
$K = 2$	188	0.37	Binary search 0.30	1	0.89
$K = 1$	181	0.39	Binary search 0.33	1	0.82
$\theta = 0.25$	226	0.31	Binary search 0.26	1	0.91
$\theta = 0.375$	177	0.38	Binary search 0.31	1	0.96
$\theta = 0.625$	116	0.50	Binary search 0.38	1	0.78
$\theta = 0.75$	107	0.54	Binary search 0.39	1	0.82
One rescuing step	102	0.33	Binary search 0.27	1	1.00
No null replays	203	0.37	Judge, thinking 0.42	6	0.60
Split samples	193	0.34	Binary search 0.27	1	0.96
Observed action as control	157 [†]	0.44	Judge, thinking 0.40	1	0.82
Second oracle (Flash)	120	0.38	Binary search 0.29	1	0.78
Both oracles averaged	121	0.43	Binary search 0.32	1	0.82
Flagged corrections dropped	152	0.43	Binary search 0.34	1	1.00
Oracle without grading	157 [†]	0.05	Counterfactual search@40 0.11	7	0.24

Table 6: Method scores under other reference choices (post hoc), pooled over the three domains: rescuable failures (re-selected under each choice; [†]: the registered 157 kept as a fixed cohort), the first write’s mean $R(\hat{k})$, the best LLM method and its $R(\hat{k})$ (judge, thinking: all-at-once with Pro and thinking), the first write’s rank among the ten methods of Table 1, and Kendall’s τ of the method ordering against the registered one ($K = 4$ oracle samples, null replays, threshold $\theta = 0.5$). Rows with fewer samples average over all subsets of the four; split samples choose rescuable failures on two samples and two null replays and score on the others; the observed action forced at step k replaces null replays and the oracle without grading replaces the registered oracle, both scored on the fixed cohort; the second oracle is V4-Flash, “both oracles averaged” averages the two oracles’ profiles, and “flagged corrections dropped” removes the corrections the admissibility screen flagged for an unseen value or two more calls.

turn, but no more than the agent’s own blocks do.

Individual profiles. Choosing the best rescue step on two oracle samples and two null replays and re-choosing it on the other two (all 12 splits, all 300 failures), the halves agreed for 56% of the failures both found rescuable (72% within one step; 52–58% by domain), and the step chosen on one half gained 0.48 on the other against 0.75 in sample. With all four samples, two independent runs

Method	Definition	Retail	Airline	AppWorld	Pooled [95%]	Exact	Position	\$	Difference [95%]
All-at-once, Flash	earliest (reg.)	0.29	0.41	0.25	0.30 [0.24, 0.37]	0.28	0.64	0.0010	
	most likely rescue	0.36	0.35	0.26	0.32 [0.25, 0.39]	0.32	0.68	0.0010	+0.02 [−0.04, +0.08]
	largest gain over null	0.29	0.41	0.28	0.31 [0.24, 0.39]	0.34	0.64	0.0010	+0.01 [−0.04, +0.06]
All-at-once, Pro	earliest (reg.)	0.30	0.34	0.26	0.30 [0.23, 0.36]	0.31	0.63	0.0031	
	most likely rescue	0.33	0.47	0.33	0.36 [0.29, 0.44]	0.38	0.65	0.0031	+0.07 [+0.01, +0.12]
	largest gain over null	0.39	0.42	0.29	0.36 [0.29, 0.44]	0.36	0.67	0.0031	+0.07 [+0.02, +0.12]
	earliest, no grading	0.17	0.32	0.25	0.23 [0.17, 0.29]	0.24	0.58	0.0028	−0.07 [−0.13, −0.00]
Binary search, Pro	earliest (reg.)	0.33	0.34	0.35	0.34 [0.27, 0.41]	0.32	0.69	0.0064	
	most likely rescue	0.38	0.42	0.32	0.37 [0.29, 0.44]	0.35	0.72	0.0069	+0.03 [−0.04, +0.10]
	largest gain over null	0.44	0.42	0.35	0.40 [0.33, 0.48]	0.39	0.74	0.0071	+0.07 [−0.01, +0.13]
	earliest, no grading	0.21	0.24	0.24	0.23 [0.16, 0.29]	0.20	0.66	0.0064	−0.11 [−0.20, −0.02]
Counterfactual search@0	first suspect, no grading	0.22	0.13	0.28	0.22 [0.16, 0.28]	0.20	0.52	0.0037	
Counterfactual search@40	no grading	0.22	0.31	0.28	0.26 [0.20, 0.33]	0.26	0.57	0.0222	+0.04 [+0.00, +0.09]
First write	rule	0.53	0.27	0.39	0.42 [0.33, 0.51]	0.38	0.78	0	

Table 7: Judges with their registered prompt (earliest decisive mistake) and post hoc variants: the step most likely to rescue the task, the step with the largest gain over letting the agent act again, and the registered definition without grading information or hidden instructions (what counterfactual search sees); and search with and without its replays. On the 157 rescuable failures: mean $R(\hat{k})$, exact agreement with the best rescue step, mean relative position of the named step, dollars per failure (mean over domains), and the paired difference in pooled $R(\hat{k})$ from the registered prompt (search@40: from search@0) with a cluster-bootstrap interval.

Method	Retail (65)	Airline (37)	τ^2 (102) [95%]	Exact	Gold deviation — method [95%]
Gold deviation (rule)	0.58	0.35	0.50 [0.38, 0.60]	0.49	
First write (rule)	0.53	0.27	0.44 [0.32, 0.55]	0.43	+0.06 [+0.01, +0.12]
All-at-once, Pro	0.30	0.34	0.32 [0.23, 0.40]	0.33	+0.18 [+0.02, +0.33]
Binary search, Pro	0.33	0.34	0.33 [0.24, 0.43]	0.34	+0.16 [+0.02, +0.30]
Counterfactual search@40	0.22	0.31	0.25 [0.17, 0.34]	0.25	+0.24 [+0.09, +0.38]

Table 8: The gold-deviation rule (post hoc): the first state-changing call that matches no unused gold action, else the last step (20% of failures). Mean $R(\hat{k})$ on the 102 rescuable τ^2 failures and paired differences with cluster-bootstrap intervals.

Method	Model	Proposed	Gain if proposed	$R(\hat{k})$
First write	rule	0.64	0.61	0.42
Last step	rule	0.53	0.33	0.20
All-at-once	Flash	0.58	0.51	0.30
All-at-once	Pro	0.60	0.48	0.30
All-at-once	Pro, thinking	0.61	0.49	0.31
Step-by-step	Pro	0.52	0.34	0.20
Binary search	Pro	0.64	0.49	0.34
Root-cause attributor	Pro	0.56	0.46	0.27
RRSI trace digester	Pro	0.57	0.34	0.22
Counterfactual search@40	Pro, replays	0.59	0.44	0.26
Best rescue step		0.88	0.87	0.76

Domain	Step	Steps	Proposed	Gain if proposed	Null
Retail	utterance	630	0.35	0.13	0.25
	read	400	0.24	0.02	0.45
	write	125	0.51	0.52	0.04
Airline	utterance	569	0.38	0.13	0.16
	read	321	0.33	0.08	0.31
	write	127	0.48	0.18	0.02
AppWorld	docs	535	0.07	0.20	0.13
	read	885	0.13	0.12	0.13
	write	272	0.53	0.21	0.13
	submit	91	0.65	0.32	0.01

Table 9: $R(\hat{k})$ split into the share of oracle samples that proposed an applicable correction at the named step and the mean gain of those corrections (over failures with at least one), on the 157 rescuable failures (post hoc).

Table 10: The same split over all steps of all 300 failures by kind of step, with the success of null replays where they were run (post hoc; τ^2 transfers and other AppWorld code omitted).

agreed for 15 of 17 failures. Profiles therefore rank methods consistently while individual best steps carry noticeable selection noise.

Oracle without grading. The oracle got the registered instructions with the grading clauses removed, the same context (policy and tools, or the agent’s rules and APIs) and the episode up to the

step, and was told that the episode failed; it ran at every step any method of Table 1 named on the rescuable failures and at their best and earliest rescuing steps, with 4 samples, 2 corrected replays each and the registered null replays (new ones where the registered oracle had never corrected the step). At best rescue steps it proposed a correction in 42% of samples, against 88% with the grading.

Domain	Corrections	Unseen	Unseen	More calls		Calls
	scored	value	number	$\geq +2$	$\geq +1$	obs. / corr.
Retail	1,554	0.4	7.6	0.4	1.5	0.94 / 0.68
Airline	1,549	1.6	10.0	2.3	3.7	1.30 / 0.80
AppWorld	1,692	0.2	1.1	2.7	15.9	0.96 / 1.04
At best rescue steps	553	0.2	2.9	2.9	5.8	

Table 11: Mechanical screen of the oracle’s corrections (post hoc): every correction that entered a registered profile, in %. Unseen value: an argument (an identifier, name, e-mail or date; in AppWorld a value passed to an API to look something up) that appears neither in the episode before the step nor in the policy, tool schemas or API documentation. Unseen number: a number above 10 likewise absent, often a sum or difference the agent could compute. More calls: tool or API calls beyond $\max(1, \text{observed})$. Re-scored without the corrections flagged for an unseen value or two more calls, the ranking was unchanged (Table 6).

Only 34 of the 157 failures had a scanned step with a gain of at least 0.5, and its best scanned step was the registered best rescue step for 29%. The first write fell below binary search asked for the largest gain ($-0.04 [-0.08, -0.003]$), the all-at-once judge so asked ($-0.05 [-0.10, -0.01]$) and informed search ($-0.05 [-0.10, -0.01]$), and did not differ detectably from registered binary search ($-0.01 [-0.04, +0.02]$). It scored near zero in retail (0.01) but still highest in AppWorld (0.10, against 0.09 for search and binary search), whose writes kept a quarter of their gain.

Fresh replays in the repair experiment. The oracle pointers lost 0.12 $[0.06, 0.17]$ of their gain on fresh replays across all 81 traces (0.39 to 0.27) but stayed above RRSI’s evidence ($+0.18 [+0.11, +0.27]$) and binary search’s ($+0.06 [-0.01, +0.14]$; on rescuable failures $+0.17 [+0.01, +0.34]$); pointers chosen without the profile moved by -0.02 to $+0.02$ on all traces. The correction given to the oracle-fix group, replayed 8 more times at its step against 8 new null replays, succeeded in 0.80 of replays against 0.95 in the profile, with null success 0.15 against 0.10: its gain fell from 0.85 to 0.65 ($-0.19 [-0.27, -0.11]$), stayed at least 0.5 on 70% of the 60 traces and above 0 on 83%.

Supplementary material. The supplement holds both registrations with their addenda and a timeline of commit dates, every replay’s outcome (failure, step, sample, kind of replay, success, whether the correction was applied), the profiles, every method’s named steps, the repair experiment’s

rounds and deployments, and the code that computes every table from them.

G Registration and Deviations

The registration (methods, prompts, metrics, the primary comparison and the budget) was committed before any method ran on the main set. Changes to the attribution study, each logged before the data it affected were analysed:

- **Reference protocol.** The registration left the protocol to be chosen on the pilot by test-retest agreement. After the earliest-step protocol agreed on 42% of pilot pairs, protocol 3 and the rescue-gain metrics replaced it (Addendum 1), while methods were running and before any main-set profile existed.
- **Retail top-up.** The first loop’s harnesses yielded 86 distinct retail failures under the sampling rule; 14 more were drawn with the same rule from two further runs of the loop, as registered. All methods moved by at most 0.02 when they were added.
- **AppWorld.** Registered in Addendum 2 before any AppWorld failure was drawn. Two corrected replays of one step failed on every attempt because the oracle’s code passed a literal ellipsis that AppWorld’s request log cannot serialise; they count as corrections not applied. That failure’s other corrections at the step rescued nothing and its largest gain was 0.375, so they could not have made it rescuable.

Repair experiment. Registered separately, before any attribution was run on its traces and before any candidate was drafted. As checks, the RRSI-analysis group’s mean gain (-0.34) matched that of the candidates the original rounds had drafted from the same evidence (-0.42), and re-deploying each incumbent moved its held-out success by -2.5 to $+3.3$ points against its earlier deployment, within the noise of 240 episodes. The first state served as a pilot; it exposed no bug, so its cells were kept, as registered. Two operational problems had no effect on data: a lookup by abbreviated commit hash re-invoked finished deployments, which skipped every episode as already done, and a container restart interrupted the deployments, which resumed with finished episodes kept. One secondary test changed. The registered code tested the round-level contrasts by permuting slots, but a round has

Method	Model	Hand-crafted (58)		Generated (126)	
		agent	step	agent	step
All-at-once	Flash	0.59	0.09	0.58	0.21
All-at-once	Pro	0.62	0.03	0.67	0.18
All-at-once	Pro, thinking	0.62	0.15	0.65	0.25
Step-by-step	Pro	0.24	0.19	0.44	0.32
Binary search	Pro	0.26	0.09	0.38	0.04

Table 12: Who&When, with the task’s correct answer shown to the judge as in its authors’ setting: accuracy of the responsible agent and step.

Evidence	Accepted	Gains of		Round	gain	– RRSI	95% interval		p
		accepted	Round				states	runs	
No attribution	2/8	−0.8, +0.0	−0.10	+0.00	[−0.37, +0.37]	[−0.54, +0.54]	1.00		
RRSI analysis	1/8	−0.8	−0.10						
First write	2/8	+2.5, +1.2	+0.47	+0.57	[−0.32, +1.46]	[−0.49, +1.63]	0.50		
Binary search	0/8	−	+0.00	+0.10	[−0.14, +0.35]	[−0.23, +0.44]	1.00		
Counterfactual search@40	2/8	+3.3, −2.1	+0.16	+0.26	[−0.99, +1.51]	[−1.52, +2.04]	0.75		
Oracle step	1/8	+2.1	+0.26	+0.36	[−0.27, +0.99]	[−0.80, +1.52]	1.00		
RRSI analysis, no reading	2/8	−3.8, +3.3	−0.05	+0.05	[−1.55, +1.66]	[−2.52, +2.63]	1.00		
Oracle fix, no reading	1/8	+0.8	+0.10	+0.21	[−0.11, +0.53]	[−0.17, +0.59]	0.50		
Any step pointer – RRSI				+0.33	[−0.13, +0.78]	[−0.39, +1.04]	0.25		
Oracle evidence – RRSI, pooled				+0.26	[−0.35, +0.87]	[−0.49, +1.01]	0.50		
Oracle fix – RRSI, no reading				+0.16	[−1.44, +1.76]	[−2.16, +2.47]	0.75		

Table 13: The accepted round in the repair experiment: rounds of 8 in which RRSI’s own rule accepted a candidate, those candidates’ held-out gains, the round’s gain (the accepted candidate’s, 0 if none) averaged over states, and its difference from RRSI’s analysis with 95% t intervals over the 8 states and over the 4 runs (each run’s two checkpoints averaged first), and the exact sign-flip p over runs. Lower block: pooled groups. Points.

Best rescue step	n	First			Registered			Post hoc		
		write	AAO	BS	CF	AAO _g	BS _g	CF _i		
Retail, utterance	28	0.24	0.36	0.31	0.25	0.47	0.41	0.39		
Retail, write	36	0.76	0.25	0.35	0.19	0.34	0.47	0.30		
Airline, utterance	25	0.23	0.35	0.36	0.40	0.40	0.38	0.43		
Airline, write	8	0.53	0.39	0.39	0.12	0.58	0.73	0.31		
AppWorld, other code	5	0.28	0.65	0.53	0.38	0.40	0.50	0.65		
AppWorld, read	9	0.28	0.10	0.18	0.42	0.11	0.18	0.35		
AppWorld, submission	9	0.56	0.14	0.28	0.30	0.27	0.33	0.21		
AppWorld, write	29	0.42	0.31	0.42	0.24	0.37	0.41	0.36		
Best step changes state	82	0.60	0.27	0.37	0.21	0.37	0.46	0.31		
Other best steps	75	0.23	0.32	0.30	0.32	0.36	0.34	0.39		
Cells weighted equally	8	0.41	0.32	0.35	0.29	0.37	0.43	0.37		

Table 14: Mean $R(\hat{k})$ on the 157 rescuable failures by the kind of their best rescue step (cells with at least 5 failures). AAO, BS: all-at-once and binary search with Pro; CF: counterfactual search@40; _g: asked for the largest gain over null; _i: given the grading information. State-changing steps are writes and AppWorld’s submission; the last row weights the eight cells above equally.

Method	Model	All three		No grading, by domain			
		Grading	None	Retail	Airline	AppW.	
First write	rule	0.42	0.05	0.01	0.07	0.10	
Last step	rule	0.20	0.04	0.02	0.01	0.07	
All-at-once	Flash	0.30	0.07	0.03	0.20	0.04	
All-at-once	Pro	0.30	0.08	0.04	0.14	0.08	
All-at-once	Pro, thinking	0.31	0.08	0.03	0.21	0.06	
Step-by-step	Pro	0.20	0.03	0.01	0.06	0.04	
Binary search	Pro	0.34	0.06	0.02	0.10	0.09	
Root-cause attributor	Pro	0.27	0.05	0.02	0.19	0.01	
RRSI trace digester	Pro	0.22	0.07	0.02	0.16	0.05	
Counterfactual search@40	Pro, replays	0.26	0.11	0.06	0.23	0.09	
All-at-once, gain	Pro	0.36	0.10	0.07	0.22	0.07	
Binary search, gain	Pro	0.40	0.09	0.06	0.15	0.09	
All-at-once, no grading	Pro	0.23	0.07	0.02	0.15	0.08	
Binary search, no grading	Pro	0.23	0.04	0.03	0.10	0.03	
CF search@0	Pro	0.22	0.08	0.07	0.08	0.08	
CF search@0, grading	Pro	0.27	0.05	0.03	0.04	0.08	
CF search@40, grading	Pro, replays	0.35	0.11	0.07	0.20	0.09	

Table 15: Mean $R(\hat{k})$ on the 157 rescuable failures under the registered oracle (grading) and under the same oracle without the grading (none), scored at the same named steps (post hoc). Lower block: the post hoc variants of Table 1.

one accepted candidate, so each cell’s value was copied to its two slots and the permutation split the copies between groups, which is not a valid test. We noticed this after the first run of the analysis, which gave $p = .009$ for first write against no attribution, and replaced the test by an exact sign-flip test over the eight state differences ($p \geq .25$);

estimates and intervals are unchanged.

Post-hoc analyses. Not registered: the re-scoring under other reference choices, the null-replay rates by step position, the decomposition of the null-free score and the retest agreement among

Pointer	All 81 traces		33 rescuable	
	Profile	Fresh	Profile	Fresh
Oracle step	0.39	0.27	0.74	0.59
First write	0.12	0.15	0.27	0.31
Binary search	0.19	0.21	0.41	0.42
CF search@40	0.13	0.12	0.29	0.25
RRSI (earliest cited)	0.10	0.09	0.26	0.16

Table 16: Rescue gain of each repair group’s pointers in the registered profile of the experiment’s failing traces and in fresh replays: 4 new oracle samples with 2 corrected replays each and 4 new null replays per step (8 at oracle steps); post hoc.

split-half agreement, the oracle without grading
and the fresh replays (attrib/review4.py,
attrib/review4_runs.py).

1195
1196
1197

failures both runs found rescuable (Table 6,
Figure 2; attrib/robustness.py); the split of
 $R(\hat{k})$ into proposals and their gains, search’s
first suspect, the null replays’ reproduction of
the observed action and the split-sample scores
(Tables 9, 10; attrib/review2.py); the observed-
action control (attrib/orig_control.py); the
second oracle (attrib/second_oracle.py,
attrib/groundtruth3.py with another oracle
model, sharing null replays where both oracles
flagged a step); the judges with other defini-
tions or without grading information (Table 7;
attrib/rescue_prompt.py); the gold-deviation
rule (Table 8; attrib/gold_rule.py); the admis-
sibility screen and the audit sample (Table 11;
attrib/admissibility.py); counterfactual
search with grading information and the factor-
ial analysis (attrib/search.py -informed,
attrib/factorial.py); in the repair experiment,
the rescue score of each group’s pointers on its
own failures (attrib/phaseb_accuracy.py),
the oracle-pointer group, run after the reg-
istered analysis with held-out results un-
seen until its deployments had finished
(attrib/phaseb_oracle.py), the proposers’
reading of traces (attrib/phaseb_reading.py),
the no-attribution group’s mean without its worst
candidate, the intervals with states grouped
by run and the correlation of pointer scores
with gains (attrib/phaseb_review3.py),
the two no-reading groups, run with RRSI’s
proposer blocked from failing traces and
held-out results unseen until their deploy-
ments had finished (attrib/phaseb.py,
attrib/phaseb_review3.py -noread); and
the checks of Appendix F: accepted-round
intervals by run, scores by kind of best rescue
step, admissibility by kind, intervention scope,